

```
In [1]: #import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

#ignore/ disable warnings
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: #Read data from csv file
df=pd.read_csv(r"C:\Users\Sanjay Lohar\Downloads\Investment.csv")
```

```
In [3]: df.describe()
```

```
Out[3]:
```

| | age | duration | campaign | pdays | previous | emp_var_rate | cons_price_idx | cons_conf_idx | euribor3m | nr |
|-------|-------------|--------------|--------------|--------------|--------------|--------------|----------------|---------------|--------------|----|
| count | 41188.00000 | 41188.000000 | 41188.000000 | 41188.000000 | 41188.000000 | 41188.000000 | 41188.000000 | 41188.000000 | 41188.000000 | 41 |
| mean | 40.02406 | 258.285010 | 2.567593 | 962.475454 | 0.172963 | 0.081886 | 93.575664 | -40.502600 | 3.621291 | 5 |
| std | 10.42125 | 259.279249 | 2.770014 | 186.910907 | 0.494901 | 1.570960 | 0.578840 | 4.628198 | 1.734447 | |
| min | 17.00000 | 0.000000 | 1.000000 | 0.000000 | 0.000000 | -3.400000 | 92.201000 | -50.800000 | 0.634000 | 4 |
| 25% | 32.00000 | 102.000000 | 1.000000 | 999.000000 | 0.000000 | -1.800000 | 93.075000 | -42.700000 | 1.344000 | 5 |
| 50% | 38.00000 | 180.000000 | 2.000000 | 999.000000 | 0.000000 | 1.100000 | 93.749000 | -41.800000 | 4.857000 | 5 |
| 75% | 47.00000 | 319.000000 | 3.000000 | 999.000000 | 0.000000 | 1.400000 | 93.994000 | -36.400000 | 4.961000 | 5 |
| max | 98.00000 | 4918.000000 | 56.000000 | 999.000000 | 7.000000 | 1.400000 | 94.767000 | -26.900000 | 5.045000 | 5 |

```
In [4]: #Target variable
df['Invested']=df['Invested'].replace(['Yes','No'],[1,0])
```

```
In [5]: df.head()
```

```
Out[5]:
```

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | ... | campaign | pdays | previous | poutcome |
|---|-----|-------------|---------|-------------------|---------|---------|------|----------|-------|-------------|-----|----------|-------|----------|----------|
| 0 | 44 | blue-collar | married | basic.4y | unknown | yes | no | cellular | aug | thu | ... | 1 | 999 | 0 | nonexi |
| 1 | 53 | technician | married | unknown | no | no | no | cellular | nov | fri | ... | 1 | 999 | 0 | nonexi |
| 2 | 28 | management | single | university.degree | no | yes | no | cellular | jun | thu | ... | 3 | 6 | 2 | suc |
| 3 | 39 | services | married | high.school | no | no | no | cellular | apr | fri | ... | 2 | 999 | 0 | nonexi |
| 4 | 55 | retired | married | basic.4y | no | yes | no | cellular | aug | fri | ... | 1 | 3 | 1 | suc |

5 rows × 21 columns

```
In [6]: df.dtypes
```

```
Out[6]:
```

| | |
|----------------|---------|
| age | int64 |
| job | object |
| marital | object |
| education | object |
| default | object |
| housing | object |
| loan | object |
| contact | object |
| month | object |
| day_of_week | object |
| duration | int64 |
| campaign | int64 |
| pdays | int64 |
| previous | int64 |
| poutcome | object |
| emp_var_rate | float64 |
| cons_price_idx | float64 |
| cons_conf_idx | float64 |
| euribor3m | float64 |
| nr_employed | float64 |
| Invested | int64 |
| dtype: | object |

```
In [7]: #check missing values
#treat missing values (if any)
df.isnull().sum()
```

```
Out[7]: age          0
        job          0
        marital      0
        education    0
        default      0
        housing      0
        loan         0
        contact      0
        month        0
        day_of_week  0
        duration     0
        campaign     0
        pdays        0
        previous     0
        poutcome     0
        emp_var_rate 0
        cons_price_idx 0
        cons_conf_idx 0
        euribor3m    0
        nr_employed  0
        Invested     0
        dtype: int64
```

```
In [8]: #Check unique entries in job column
df['job'].unique()
```

```
Out[8]: array(['blue-collar', 'technician', 'management', 'services', 'retired',
              'admin.', 'housemaid', 'unemployed', 'entrepreneur',
              'self-employed', 'unknown', 'student'], dtype=object)
```

```
In [9]: #Check marital status
df['marital'].unique()
```

```
Out[9]: array(['married', 'single', 'divorced', 'unknown'], dtype=object)
```

```
In [10]: #Check unique entries in education column
df['education'].unique()
```

```
Out[10]: array(['basic.4y', 'unknown', 'university.degree', 'high.school',
              'basic.9y', 'professional.course', 'basic.6y', 'illiterate'],
              dtype=object)
```

```
In [11]: #Replace
#basic.4y. basic.6y,basic.9y with basic
df['education']=df['education'].replace(['basic.4y','basic.6y','basic.9y'],
                                       ['basic','basic','basic'])
```

```
In [12]: df['education'].unique()
```

```
Out[12]: array(['basic', 'unknown', 'university.degree', 'high.school',
              'professional.course', 'illiterate'], dtype=object)
```

```
In [13]: df['default'].unique()
```

```
Out[13]: array(['unknown', 'no', 'yes'], dtype=object)
```

```
In [14]: df['housing'].unique()
```

```
Out[14]: array(['yes', 'no', 'unknown'], dtype=object)
```

```
In [15]: df['loan'].unique()
```

```
Out[15]: array(['no', 'yes', 'unknown'], dtype=object)
```

```
In [16]: df['loan'].unique()
```

```
Out[16]: array(['no', 'yes', 'unknown'], dtype=object)
```

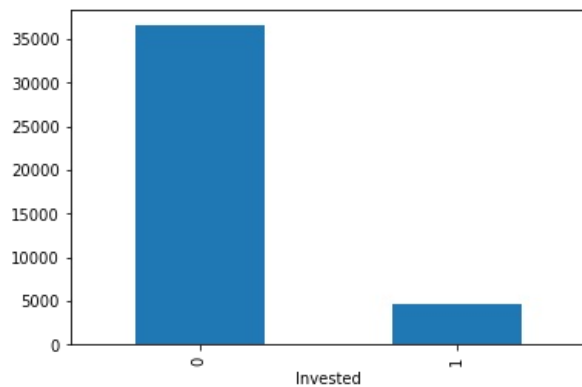
```
In [17]: df['poutcome'].unique()
```

```
Out[17]: array(['nonexistent', 'success', 'failure'], dtype=object)
```

```
In [18]: #----- Explore Data (EDA) -----
```

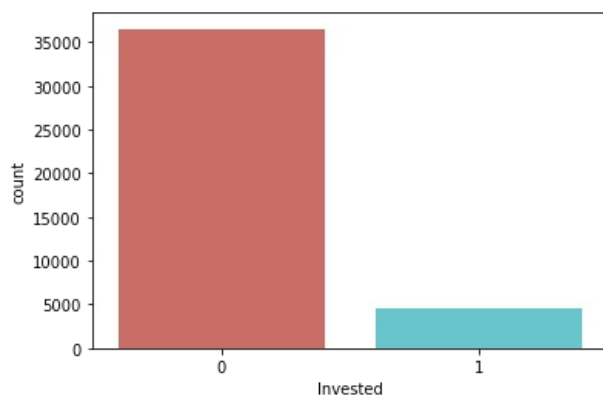
```
In [19]: df.groupby('Invested')['Invested'].count().plot(kind='bar')
```

```
Out[19]: <AxesSubplot:xlabel='Invested'>
```



In [20]: `sns.countplot(x='Invested', data = df, palette='hls') #hue, saturation, and luminance for automatic color highl`

Out[20]: `<AxesSubplot:xlabel='Invested', ylabel='count'>`



In [21]: `df.groupby('Invested')['Invested'].count()`

Out[21]:

| Invested | count |
|----------|-------|
| 0 | 36548 |
| 1 | 4640 |


Name: Invested, dtype: int64

In [22]: `#Our classes are imbalanced. Only 11% of the customers have invested.
4640/(4640+36548)`

Out[22]: `0.11265417111780131`


In [23]: `df['age'].hist()`

Out[23]: `<AxesSubplot:>`



In [24]: `sns.distplot(df['age'])`

Out[24]: `<AxesSubplot:xlabel='age', ylabel='Density'>`




```
In [25]: df.groupby('Invested').mean()
```

```
Out[25]:
```


| | age | duration | campaign | pdays | previous | emp_var_rate | cons_price_idx | cons_conf_idx | euribor3m | nr_employed |
|----------|-----------|------------|----------|------------|----------|--------------|----------------|---------------|-----------|-------------|
| Invested | | | | | | | | | | |
| 0 | 39.911185 | 220.844807 | 2.633085 | 984.113878 | 0.132374 | 0.248875 | 93.603757 | -40.593097 | 3.811491 | 5176.166600 |
| 1 | 40.913147 | 553.191164 | 2.051724 | 792.035560 | 0.492672 | -1.233448 | 93.354386 | -39.789784 | 2.123135 | 5095.115991 |

```
In [26]: pd.crosstab(df['job'], df['Invested']).plot(kind='bar', stacked=True)
plt.xlabel('Job Type')
plt.ylabel('count of cust')
plt.show()
```




```
In [27]: pd.crosstab(df['marital'], df['Invested']).plot(kind='bar', stacked=True)
```

```
Out[27]: <AxesSubplot: xlabel='marital'>
```



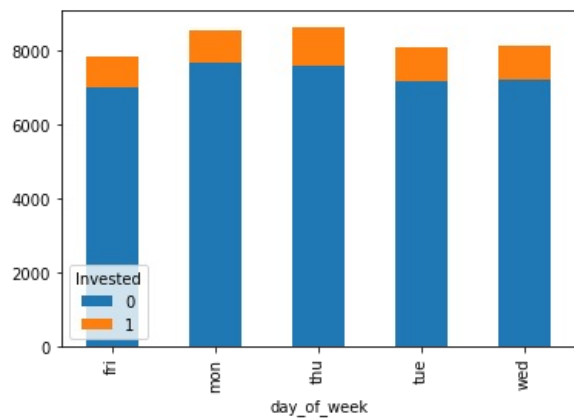
```
In [28]: pd.crosstab(df['education'], df['Invested']).plot(kind='bar', stacked=True)
```

```
Out[28]: <AxesSubplot: xlabel='education'>
```




```
In [29]: pd.crosstab(df['day_of_week'], df['Invested']).plot(kind='bar', stacked=True)
```

```
Out[29]: <AxesSubplot:xlabel='day_of_week'>
```




```
In [30]: pd.crosstab(df['month'], df['Invested']).plot(kind='bar', stacked=True)
```

```
Out[30]: <AxesSubplot:xlabel='month'>
```




```
In [31]: #----- Handle outliers -----
```


```
In [32]: #Check outliers
plt.boxplot(df['age']) #has outlier
plt.show()
```



```
In [33]: #Check outliers
plt.boxplot(df['duration']) #has outlier
plt.show()
```



```
In [34]: #Check outliers
plt.boxplot(df['pdays']) #has outlier
plt.show()
```




```
In [40]: #remove outliers
#user defined function for outlier treatment
def rm_out(d,c):
    #find q1 and q3
    q1=d[c].quantile(0.25)
    q3=d[c].quantile(0.75)

    #find iqr
    iqr=q3-q1


    #upper bound and lower bound
    ub=q3+1.5*iqr
    lb=q1-1.5*iqr

    final_data=d.loc[(d[c]>lb) & (d[c]<ub)]
    return final_data
```

```
In [83]: #Outlier treatment: age column
df=rm_out(df,'age')
plt.boxplot(df['age'])
plt.show()
```



```
In [90]: #Outlier treatment: duration column
df=rm_out(df,'duration')
plt.boxplot(df['duration'])
plt.show()
```



```
In [91]: import numpy as np
#Replace pdays 999 with NaN (missing values)
df['pdays']=df['pdays'].replace(999,np.nan)
```

```
In [92]: #Replace pdays NaN (missing values) with median
df['pdays']=df['pdays'].fillna(df['pdays'].median())
```

```
In [93]: # One-hot encoding (dummy conversion)
df_categorical = df.select_dtypes(include=['object'])
df_categorical
```

Out[93]:

| | job | marital | education | default | housing | loan | contact | month | day_of_week | poutcome |
|-------|-------------|---------|---------------------|---------|---------|------|-----------|-------|-------------|-------------|
| 0 | blue-collar | married | basic | unknown | yes | no | cellular | aug | thu | nonexistent |
| 1 | technician | married | unknown | no | no | no | cellular | nov | fri | nonexistent |
| 2 | management | single | university.degree | no | yes | no | cellular | jun | thu | success |
| 3 | services | married | high.school | no | no | no | cellular | apr | fri | nonexistent |
| 4 | retired | married | basic | no | yes | no | cellular | aug | fri | success |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 41183 | retired | married | high.school | unknown | no | yes | telephone | jun | thu | nonexistent |
| 41184 | housemaid | married | basic | unknown | no | no | telephone | may | thu | nonexistent |
| 41185 | admin. | single | university.degree | unknown | yes | yes | telephone | may | wed | nonexistent |
| 41186 | technician | married | professional.course | no | no | yes | telephone | oct | tue | nonexistent |
| 41187 | student | single | high.school | no | no | no | telephone | may | fri | nonexistent |

35688 rows × 10 columns

```
In [94]: # convert into dummies
df_dummies = pd.get_dummies(df_categorical)
df_dummies.head()
```

```
Out[94]:
```

| | job_admin. | job_blue-collar | job_entrepreneur | job_housemaid | job_management | job_retired | job_self-employed | job_services | job_student | job_technician |
|---|------------|-----------------|------------------|---------------|----------------|-------------|-------------------|--------------|-------------|----------------|
| 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
| 2 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| 3 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| 4 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |

5 rows × 11 columns

```
In [95]: #Check the correlation of numeric variables
df_numeric = df.select_dtypes(include=['float64', 'int64'])
df_numeric.head()
```

```
Out[95]:
```

| | age | duration | campaign | pdays | previous | emp_var_rate | cons_price_idx | cons_conf_idx | euribor3m | nr_employed | Invested |
|---|-----|----------|----------|-------|----------|--------------|----------------|---------------|-----------|-------------|----------|
| 0 | 44 | 210 | 1 | 6.0 | 0 | 1.4 | 93.444 | -36.1 | 4.963 | 5228.1 | 0 |
| 1 | 53 | 138 | 1 | 6.0 | 0 | -0.1 | 93.200 | -42.0 | 4.021 | 5195.8 | 0 |
| 2 | 28 | 339 | 3 | 6.0 | 2 | -1.7 | 94.055 | -39.8 | 0.729 | 4991.6 | 1 |
| 3 | 39 | 185 | 2 | 6.0 | 0 | -1.8 | 93.075 | -47.1 | 1.405 | 5099.1 | 0 |
| 4 | 55 | 137 | 1 | 3.0 | 1 | -2.9 | 92.201 | -31.4 | 0.869 | 5076.2 | 1 |

```
In [ ]:
```

```
In [96]: #combine numeric columns and dummies to create master data
master=pd.concat([df_numeric,df_dummies], axis=1)
master.head()
```

```
Out[96]:
```

| | age | duration | campaign | pdays | previous | emp_var_rate | cons_price_idx | cons_conf_idx | euribor3m | nr_employed | ... | month_oct | month_nov |
|---|-----|----------|----------|-------|----------|--------------|----------------|---------------|-----------|-------------|-----|-----------|-----------|
| 0 | 44 | 210 | 1 | 6.0 | 0 | 1.4 | 93.444 | -36.1 | 4.963 | 5228.1 | ... | 0 | 0 |
| 1 | 53 | 138 | 1 | 6.0 | 0 | -0.1 | 93.200 | -42.0 | 4.021 | 5195.8 | ... | 0 | 0 |
| 2 | 28 | 339 | 3 | 6.0 | 2 | -1.7 | 94.055 | -39.8 | 0.729 | 4991.6 | ... | 0 | 0 |
| 3 | 39 | 185 | 2 | 6.0 | 0 | -1.8 | 93.075 | -47.1 | 1.405 | 5099.1 | ... | 0 | 0 |
| 4 | 55 | 137 | 1 | 3.0 | 1 | -2.9 | 92.201 | -31.4 | 0.869 | 5076.2 | ... | 0 | 0 |

5 rows × 14 columns

```
In [97]: # Create X and Y
y=master['Invested']
x=master.drop('Invested',axis=1)
```

```
In [98]: # split data into train and test
from sklearn.model_selection import train_test_split

xtrain,xtest, ytrain, ytest = train_test_split(x, y, test_size = 0.3,
                                              random_state=0)
```

```
In [99]: print(xtrain.shape, ytrain.shape, xtest.shape, ytest.shape )
(24981, 61) (24981,) (10707, 61) (10707,)
```

```
In [100]: #----- Feature Selection / Dimensionality reduction -----
#Chi2 Test ----
```

```
In [101]: from sklearn.feature_selection import RFE
```

- We use chi2 (Chi-square) test to find the relation between categorical variables.
- The RFE (Recursive Feature Elimination) function generates the p-value

-
- H0 : correlation r = 0 # There is no correlation
 - H1 : correlation r != 0 # There is correlation
 - p < 0.05 — this means the two categorical variables are correlated. # here H0 rejected & H1 is accepted cause p < 0.05
 - p > 0.05 — this means the two categorical variables are not correlated. # here H0 accepted & H1 is rejected cause p > 0.05

```
In [102]: from sklearn.linear_model import LogisticRegression
```

```
In [106]: logarea = LogisticRegression()
```



```
logreg = LogisticRegression()
rfe = RFE(logreg, n_features_to_select=20)
rfe = rfe.fit(xtrain, ytrain)
print(rfe.support_)
#print(rfe.ranking_)
```

```
[False False False False True True False False False False False True
 False False False True False True True False False False False False
 False False False False False False False False True True False False
 True False False False False False True False True False False False
 True True True True True False True False False False True True
 True]
```

```
#print the significant variable names
xtrain.columns[rfe.support_]
```

```
Index(['previous', 'emp_var_rate', 'job_blue-collar', 'job_retired',
       'job_services', 'job_student', 'default_no', 'default_unknown',
       'housing_unknown', 'contact_telephone', 'month_aug', 'month_mar',
       'month_may', 'month_nov', 'month_oct', 'month_sep', 'day_of_week_mon',
       'poutcome_failure', 'poutcome_nonexistent', 'poutcome_success'],
      dtype='object')
```

```
##store significant variables
xtrain_new=xtrain[xtrain.columns[rfe.support_]]
xtrain_new
```

| | previous | emp_var_rate | job_blue-collar | job_retired | job_services | job_student | default_no | default_unknown | housing_unknown | contact_te |
|-------|----------|--------------|-----------------|-------------|--------------|-------------|------------|-----------------|-----------------|------------|
| 30987 | 1 | -1.8 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | |
| 23281 | 0 | 1.4 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | |
| 18754 | 0 | 1.4 | 0 | 0 | 1 | 0 | 0 | 1 | 0 | |
| 835 | 0 | 1.4 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 17547 | 0 | -0.1 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 23932 | 0 | -1.8 | 0 | 1 | 0 | 0 | 1 | 0 | 0 | |
| 37052 | 0 | -0.1 | 1 | 0 | 0 | 0 | 0 | 1 | 0 | |
| 35120 | 1 | -2.9 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 24493 | 0 | 1.1 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |
| 3118 | 0 | -1.8 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | |

24981 rows × 20 columns

```
#-----End of feature selection -----
```

Model1: Logistic Regression

```
from sklearn.linear_model import LogisticRegression
```

```
#create model object
logreg=LogisticRegression()
```

```
#Fit the model using training data
logreg.fit(xtrain,ytrain)
```

```
#check the accuracy of training model
logreg.score(xtrain,ytrain)
```

```
0.946919658940795
```

```
#predict y using test data
y_pred=logreg.predict(xtest)
```

```
#check the accuracy of test model
logreg.score(xtest,ytest)
```

```
0.9522742131315961
```

----- Measure to test the model -----

- Accuracy = (TP+TN)/(TP+TN+FP+FN)
- Precision=TP/(TP+FP)
- Recall= TP/(TP+FN) ... also called sensitivity/ hit rate / True Positive Rate(TPR)
- F1 Score= 2 x [(precision * recall)/(precision+recall)]

- F1 score calculates the harmonic mean between precision and recall.
- It generates a score between 0 (being lowest) and 1 (being highest)
-
- Specificity = $TN / (TN + FP)$...also called selectivity / True Negative Rate (TNR)

```
In [112]: from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score, confusion_matrix
acc = accuracy_score(ytest, y_pred)
prec = precision_score(ytest, y_pred)
rec = recall_score(ytest, y_pred)
f1 = f1_score(ytest, y_pred)

print('Accuracy:', acc, '\nPrecision:', prec, '\nRecall:', rec, '\nF1 Score:', f1)

results = pd.DataFrame([['Logistic Regression', acc, prec, rec, f1]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
results

Accuracy: 0.9522742131315961
Precision: 0.6965699208443272
Recall: 0.4
F1 Score: 0.5081809432146295
```

```
Out[112]:
```

| | Model | Accuracy | Precision | Recall | F1 Score |
|---|---------------------|----------|-----------|--------|----------|
| 0 | Logistic Regression | 0.952274 | 0.69657 | 0.4 | 0.508181 |

```
In [113]: prec = precision_score(ytest, y_pred)
prec
```

```
Out[113]: 0.6965699208443272
```

```
In [114]: #import confusion matrix library
from sklearn.metrics import confusion_matrix
```

```
In [115]: #create confusion matrix
confusion_matrix(ytest, y_pred)
```

```
Out[115]: array([[9932, 115],
                [ 396, 264]], dtype=int64)
```

```
In [116]: #Precision: Out of all predicted positives, how many are really positive
#Precision = TP / (TP + FP)
264 / (264 + 115)
```

```
Out[116]: 0.6965699208443272
```

```
In [117]: #Accuracy = (TP + TN) / (TP + TN + FN + FP)
(9932 + 264) / (9932 + 264 + 396 + 115)
```

```
Out[117]: 0.9522742131315961
```

```
In [118]: #TPR or Sensitivity or recall TPR = TP / TP + FN
264 / (396 + 264)
```

```
Out[118]: 0.4
```

```
In [119]: #TNR or Specificity or selectivity TNR = TN / TN + FP
9932 / (9932 + 115)
```

```
Out[119]: 0.9885537971533791
```

```
In [ ]:
```

-----KNN-----

```
In [120]: #import knn library from sklearn
from sklearn.neighbors import KNeighborsClassifier
```

```
In [121]: #Create model object
knn = KNeighborsClassifier(n_neighbors=5, metric='euclidean')
```

```
In [122]: #fit the training model
knn.fit(xtrain, ytrain)
```

```
Out[122]: KNeighborsClassifier(metric='euclidean')
```

```
In [123]: #test the accuracy of training model
knn.score(xtrain, ytrain)
```

Out[123]: 0.9558064128737841

```
In [124... #test the model using xtest
y_pred=knn.predict(xtest)
```

```
In [125... #Check accuracy of test model
knn.score(xtest,ytest)
```

Out[125]: 0.9463902120108341

```
In [126... #import confusion matrix from sklearn
from sklearn.metrics import confusion_matrix
```

```
In [127... confusion_matrix(ytest,y_pred)
```

Out[127]: array([[9854, 193],
 [381, 279]], dtype=int64)

```
In [128... #Calculate Recall, Precision, Sensitivity
knn_acc=(9854+279)/(9854+279+193+381)
knn_rec=279/(279+381)
knn_prec=(279/(279+193))
knn_spec=9854/(9854+193)
knn_f1=2*(knn_prec*knn_rec)/(knn_prec + knn_rec)
```

```
In [129... results = pd.DataFrame([
                                ['Log Reg', acc,prec,rec,f1],
                                ['KNN', knn_acc,knn_prec,knn_rec,knn_f1]
                                ],
                                columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results
```

Out[129]:

| | Model | Accuracy | Precision | Recall | F1 Score |
|---|---------|----------|-----------|----------|----------|
| 0 | Log Reg | 0.952274 | 0.696570 | 0.400000 | 0.508181 |
| 1 | KNN | 0.946390 | 0.591102 | 0.422727 | 0.492933 |

In []:

```
In [130... # predict probabilities
pred_prob1 = logreg.predict_proba(xtest)
pred_prob2 = knn.predict_proba(xtest)
```

```
In [131... from sklearn.metrics import roc_curve

# roc curve for models
fpr1, tpr1, thresh1 = roc_curve(ytest, pred_prob1[:,1], pos_label=1)
fpr2, tpr2, thresh2 = roc_curve(ytest, pred_prob2[:,1], pos_label=1)

# roc curve for tpr = fpr
random_probs = [0 for i in range(len(ytest))]
p_fpr, p_tpr, _ = roc_curve(ytest, random_probs, pos_label=1)
```

```
In [132... from sklearn.metrics import roc_auc_score
```

```
# auc scores
auc_score1 = roc_auc_score(ytest, pred_prob1[:,1])
auc_score2 = roc_auc_score(ytest, pred_prob2[:,1])


print(auc_score1, auc_score2)
```

0.9383393505071617 0.860920114854125

```
In [133... # matplotlib
import matplotlib.pyplot as plt
plt.style.use('seaborn')

# plot roc curves
plt.plot(fpr1, tpr1, linestyle='--',color='orange', label='Logistic Regression')
plt.plot(fpr2, tpr2, linestyle='--',color='green', label='KNN')
plt.plot(p_fpr, p_tpr, linestyle='--', color='blue')
# title
plt.title('ROC curve')
# x label
plt.xlabel('False Positive Rate')
# y label
plt.ylabel('True Positive rate')

plt.legend(loc='best')
plt.savefig('ROC')
plt.show()
```



ROC and AUC

- The Receiver Operator Characteristic (ROC) curve is an evaluation metric for binary classification problems.
- It is a probability curve that plots the TPR against FPR
- The Area Under the Curve (AUC) is the measure of the ability of a classifier to distinguish between classes and is used as a summary of the ROC curve.
- The higher the AUC, the better the performance of the model at distinguishing between the positive and negative classes.

----- Naive Bayes Classifier -----

```
In [134.. #Import Gaussian Naive Bayes model
from sklearn.naive_bayes import GaussianNB

#Create a Gaussian Classifier
gnb = GaussianNB()

# Train the model using the training sets
gnb.fit(xtrain,ytrain)

#Predict Output
y_pred = gnb.predict(xtest)
```

```
In [135.. from sklearn.metrics import confusion_matrix

confusion_matrix(ytest,y_pred)
```

```
Out[135]: array([[8967, 1080],
        [ 165,  495]], dtype=int64)
```

```
In [136.. from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score, confusion_matrix
nb_acc = accuracy_score(ytest, y_pred)
nb_prec = precision_score(ytest, y_pred)
nb_rec = recall_score(ytest, y_pred)
nb_f1 = f1_score(ytest, y_pred)

#print('Accuracy:', nb_acc, '\nPrecision:', nb_prec, '\nRecall:', nb_rec, '\nF1 Score:', nb_f1)

results = pd.DataFrame([['Log Reg', acc, prec, rec, f1],
                        ['KNN', knn_acc, knn_prec, knn_rec, knn_f1],
                        ['Naive Bayes', nb_acc, nb_prec, nb_rec, nb_f1]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
results
```

```
Out[136]:
```

| | Model | Accuracy | Precision | Recall | F1 Score |
|---|-------------|----------|-----------|----------|----------|
| 0 | Log Reg | 0.952274 | 0.696570 | 0.400000 | 0.508181 |
| 1 | KNN | 0.946390 | 0.591102 | 0.422727 | 0.492933 |
| 2 | Naive Bayes | 0.883721 | 0.314286 | 0.750000 | 0.442953 |

```
In [ ]:
```

```
In [ ]:
```

----- Decision Tree -----

```
In [137]... #import Decision Tree classifier library from sklearn
from sklearn.tree import DecisionTreeClassifier
```

```
In [138]... #Create model object
dtree=DecisionTreeClassifier(max_depth=5)
```

```
In [139]... #fit the training model
dtree.fit(xtrain,ytrain)
```

```
Out[139]: DecisionTreeClassifier(max_depth=5)
```

```
In [140]... #test the accuracy of training model
dtree.score(xtrain,ytrain)
```

```
Out[140]: 0.9526039790240582
```

```
In [141]... #test the model using xtest
y_pred=dtree.predict(xtest)
```

```
In [142]... #Check accuracy of test model
dtree.score(xtest,ytest)
```

```
Out[142]: 0.9540487531521434
```

```
In [143]... from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score, confusion_matrix
dt_acc = accuracy_score(ytest, y_pred)
dt_prec = precision_score(ytest, y_pred)
dt_rec = recall_score(ytest, y_pred)
dt_f1 = f1_score(ytest, y_pred)

results = pd.DataFrame([[ 'Log Reg', acc, prec, rec, f1],
                        [ 'KNN', knn_acc, knn_prec, knn_rec, knn_f1],
                        [ 'Naive Bayes', nb_acc, nb_prec, nb_rec, nb_f1],
                        [ 'D-Tree', dt_acc, dt_prec, dt_rec, dt_f1]],
                        columns=[ 'Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
results
```

```
Out[143]:
```

| | Model | Accuracy | Precision | Recall | F1 Score |
|---|-------------|----------|-----------|----------|----------|
| 0 | Log Reg | 0.952274 | 0.696570 | 0.400000 | 0.508181 |
| 1 | KNN | 0.946390 | 0.591102 | 0.422727 | 0.492933 |
| 2 | Naive Bayes | 0.883721 | 0.314286 | 0.750000 | 0.442953 |
| 3 | D-Tree | 0.954049 | 0.632911 | 0.606061 | 0.619195 |

----- Random Forest -----

```
In [144]... from sklearn.ensemble import RandomForestClassifier
```

- random_state– controls randomness of the sample.
- n_jobs– it tells the engine how many processors it is allowed to use. If the value is 1, it can use only one processor but if the value is -1 there is no limit.
- n_estimators– number of trees the algorithm builds before averaging the predictions.
- oob_score – OOB means out of the bag.
 - It is a random forest cross-validation method.
 - In this one-third of the sample is not used to train the data instead used to evaluate its performance.
 - These samples are called out of bag samples.

```
In [145]... rf = RandomForestClassifier(random_state=0, n_jobs=-1, max_depth=5,
                                     n_estimators=100, oob_score=True)
```

```
In [146]... #fit the model
rf.fit(xtrain, ytrain)
```

```
Out[146]: RandomForestClassifier(max_depth=5, n_jobs=-1, oob_score=True, random_state=0)
```

```
In [147]... #check OOB score (test score based on out of bag sample)
rf.oob_score_
```

```
Out[147]: 0.9450782594772027
```

```
In [148.. y_pred=rf.predict(xtest)
```

```
In [149.. rf.score(xtest, ytest)
```

```
Out[149]: 0.9490987204632484
```

```
In [150.. from sklearn.metrics import accuracy_score, f1_score,recall_score,precision_score, confusion_matrix
rf_acc = accuracy_score(ytest, y_pred)
rf_prec = precision_score(ytest, y_pred)
rf_rec = recall_score(ytest, y_pred)
rf_f1 = f1_score(ytest,y_pred)

results = pd.DataFrame(['Log Reg', acc,prec,rec,f1],
                        ['KNN', knn_acc,knn_prec,knn_rec,knn_f1],
                        ['Naive Bayes', nb_acc,nb_prec,nb_rec,nb_f1],
                        ['D-Tree', dt_acc,dt_prec,dt_rec,dt_f1],
                        ['Randm Forest', rf_acc,rf_prec,rf_rec,rf_f1]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results
```

```
Out[150]:
```

| | Model | Accuracy | Precision | Recall | F1 Score |
|---|--------------|----------|-----------|----------|----------|
| 0 | Log Reg | 0.952274 | 0.696570 | 0.400000 | 0.508181 |
| 1 | KNN | 0.946390 | 0.591102 | 0.422727 | 0.492933 |
| 2 | Naive Bayes | 0.883721 | 0.314286 | 0.750000 | 0.442953 |
| 3 | D-Tree | 0.954049 | 0.632911 | 0.606061 | 0.619195 |
| 4 | Randm Forest | 0.949099 | 0.783251 | 0.240909 | 0.368482 |

----- Support Vector Machine (SVM) -----

```
In [151.. #Support Vector Machine (SVC: Support Vector Classifier)
from sklearn.svm import SVC
svm = SVC(kernel='linear')
```

```
In [152.. #Fit the model
svm.fit(xtrain,ytrain)
```

```
Out[152]: SVC(kernel='linear')
```

```
In [153.. #check accuracy of training model
svm.score(xtrain,ytrain)
```

```
Out[153]: 0.9469596893639166
```

```
In [154.. #Predict the response from xtest
y_pred=svm.predict(xtest)
```

```
In [155.. #Check accuracy of test model
svm.score(xtest,ytest)
```

```
Out[155]: 0.9490987204632484
```

```
In [156.. from sklearn.metrics import accuracy_score, f1_score,recall_score,precision_score, confusion_matrix
svm_acc = accuracy_score(ytest, y_pred)
svm_prec = precision_score(ytest, y_pred)
svm_rec = recall_score(ytest, y_pred)
svm_f1 = f1_score(ytest,y_pred)

results = pd.DataFrame(['Log Reg', acc,prec,rec,f1],
                        ['KNN', knn_acc,knn_prec,knn_rec,knn_f1],
                        ['Naive Bayes', nb_acc,nb_prec,nb_rec,nb_f1],
                        ['D-Tree', dt_acc,dt_prec,dt_rec,dt_f1],
                        ['Randm Forest', rf_acc,rf_prec,rf_rec,rf_f1],
                        ['SVM', svm_acc,svm_prec,svm_rec,svm_f1]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results
```

```
Out[156]:
```

| | Model | Accuracy | Precision | Recall | F1 Score |
|---|--------------|----------|-----------|----------|----------|
| 0 | Log Reg | 0.952274 | 0.696570 | 0.400000 | 0.508181 |
| 1 | KNN | 0.946390 | 0.591102 | 0.422727 | 0.492933 |
| 2 | Naive Bayes | 0.883721 | 0.314286 | 0.750000 | 0.442953 |
| 3 | D-Tree | 0.954049 | 0.632911 | 0.606061 | 0.619195 |
| 4 | Randm Forest | 0.949099 | 0.783251 | 0.240909 | 0.368482 |
| 5 | SVM | 0.949099 | 0.612967 | 0.472727 | 0.533790 |

```

In [ ]:

In [157]: """ from xgboost import XGBClassifier
xgb = XGBClassifier() """

Out[157]: ' from xgboost import XGBClassifier\nxgb = XGBClassifier() '

In [158]: """ #Fit the model
xgb.fit(xtrain,ytrain) """

Out[158]: ' #Fit the model\nxgb.fit(xtrain,ytrain) '

In [159]: """ #check accuracy of training model
xgb.score(xtrain,ytrain) """

Out[159]: ' #check accuracy of training model\nxgb.score(xtrain,ytrain) '

In [160]: """ #Predict the response from xtest
y_pred=xgb.predict(xtest) """

Out[160]: ' #Predict the response from xtest\ny_pred=xgb.predict(xtest) '

In [161]: """ #Check accuracy of test model
xgb.score(xtest,ytest) """

Out[161]: ' #Check accuracy of test model\nxgb.score(xtest,ytest) '

In [162]: """ from sklearn.metrics import accuracy_score, f1_score,recall_score,precision_score, confusion_matrix
xgb_acc = accuracy_score(ytest, y_pred)
xgb_prec = precision_score(ytest, y_pred)
xgb_rec = recall_score(ytest, y_pred)
xgb_f1 = f1_score(ytest,y_pred)

results = pd.DataFrame([['Log Reg', acc,prec,rec,f1],
                        ['KNN', knn_acc,knn_prec,knn_rec,knn_f1],
                        ['Naive Bayes', nb_acc,nb_prec,nb_rec,nb_f1],
                        ['D-Tree', dt_acc,dt_prec,dt_rec,dt_f1],
                        ['Randm Forest', rf_acc,rf_prec,rf_rec,rf_f1],
                        ['SVM', svm_acc,svm_prec,svm_rec,svm_f1]
                        ['XGboost', xgb_acc, xgb_prec, xgb_rec, xgb_f1 ]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results """

Out[162]: " from sklearn.metrics import accuracy_score, f1_score,recall_score,precision_score, confusion_matrix\nxgb_acc = accuracy_score(ytest, y_pred)\nxgb_prec = precision_score(ytest, y_pred)\nxgb_rec = recall_score(ytest, y_pred)\nxgb_f1 = f1_score(ytest,y_pred)\n\nresults = pd.DataFrame([['Log Reg', acc,prec,rec,f1],\n                        ['KNN', knn_acc,knn_prec,knn_rec,knn_f1],\n                        ['Naive Bayes', nb_acc,nb_prec,nb_rec,nb_f1],\n                        ['D-Tree', dt_acc,dt_prec,dt_rec,dt_f1],\n                        ['Randm Forest', rf_acc,rf_prec,rf_rec,rf_f1],\n                        ['SVM', svm_acc,svm_prec,svm_rec,svm_f1]\n                        ['XGboost', xgb_acc, xgb_prec, xgb_rec, xgb_f1 ]],\n                        columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])\nresults "

In [ ]:

In [ ]:

In [163]: # classification report for SVM model ytest is same for all but y_pred is prediction vary or differnt by models
# homework compute classification report for all other model
# procedure is same

from sklearn.metrics import classification_report

clsreport= classification_report(ytest, y_pred)
print(clsreport)

              precision    recall  f1-score   support

    0       0.97       0.98       0.97       10047
    1       0.61       0.47       0.53         660

 accuracy
macro avg       0.79       0.73       0.75       10707
weighted avg       0.94       0.95       0.95       10707

In [ ]:

In [ ]:

In [ ]:

In [ ]:

```

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js