

# Classification algorithm project

## Algorithms in this notebook

1) Logistic Regression 2) KNN (K-nearest neighbors) 3) Decision Tree 4) Random Forest 5) SVM (Support Vector Machine) 6) Naive Bayes

```
In [1]: #import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: #Read data from csv file
df=pd.read_csv(r"C:\Users\Sanjay Lohar\Downloads\Investment.csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	campaign	pdays	previous	poutcome
0	44	blue-collar	married	basic.4y	unknown	yes	no	cellular	aug	thu ...		1	999	0	nonexi
1	53	technician	married	unknown	no	no	no	cellular	nov	fri ...		1	999	0	nonexi
2	28	management	single	university.degree	no	yes	no	cellular	jun	thu ...		3	6	2	suc
3	39	services	married	high.school	no	no	no	cellular	apr	fri ...		2	999	0	nonexi
4	55	retired	married	basic.4y	no	yes	no	cellular	aug	fri ...		1	3	1	suc

5 rows × 21 columns

```
In [4]: #Target variable
df['Invested']=df['Invested'].replace(['Yes','No'],[1,0])
```

```
In [5]: df.dtypes
```

```
Out[5]:
```

age	int64
job	object
marital	object
education	object
default	object
housing	object
loan	object
contact	object
month	object
day_of_week	object
duration	int64
campaign	int64
pdays	int64
previous	int64
poutcome	object
emp_var_rate	float64
cons_price_idx	float64
cons_conf_idx	float64
euribor3m	float64
nr_employed	float64
Invested	int64
dtype:	object

```
In [6]: #check missing values
#treat missing values (if any)
df.isnull().sum()
```

```
Out[6]: age                0
job                0
marital            0
education          0
default            0
housing            0
loan               0
contact            0
month              0
day_of_week        0
duration           0
campaign           0
pdays             0
previous           0
poutcome           0
emp_var_rate       0
cons_price_idx     0
cons_conf_idx      0
euribor3m          0
nr_employed        0
Invested           0
dtype: int64
```

```
In [7]: #print column names and data types
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 41188 entries, 0 to 41187
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   age                   41188 non-null  int64
1   job                   41188 non-null  object
2   marital               41188 non-null  object
3   education              41188 non-null  object
4   default                41188 non-null  object
5   housing                41188 non-null  object
6   loan                   41188 non-null  object
7   contact                41188 non-null  object
8   month                  41188 non-null  object
9   day_of_week            41188 non-null  object
10  duration                41188 non-null  int64
11  campaign                41188 non-null  int64
12  pdays                   41188 non-null  int64
13  previous                41188 non-null  int64
14  poutcome                41188 non-null  object
15  emp_var_rate            41188 non-null  float64
16  cons_price_idx          41188 non-null  float64
17  cons_conf_idx           41188 non-null  float64
18  euribor3m               41188 non-null  float64
19  nr_employed             41188 non-null  float64
20  Invested                41188 non-null  int64
dtypes: float64(5), int64(6), object(10)
memory usage: 6.6+ MB
```

```
In [8]: #check unique entries in job column
df['job'].unique()
```

```
Out[8]: array(['blue-collar', 'technician', 'management', 'services', 'retired',
              'admin.', 'housemaid', 'unemployed', 'entrepreneur',
              'self-employed', 'unknown', 'student'], dtype=object)
```

```
In [9]: #Replace
# technician: blue-collar
# management: white-collar
# services : white-collar
# admin. : white-collar
# housemaid : blue-collar
# entrepreneur : white-collar
# self-employed: white-collar
# unknown : unemployed
# student : unemployed
# Retired : Retired
```

```
In [10]: df['job']=df['job'].replace(['technician','housemaid',
                                     'management','services','admin.',
                                     'entrepreneur','self-employed',
                                     'unknown','student'],
                                     ['blue-collar','blue-collar',
                                     'white-collar','white-collar',
                                     'white-collar','white-collar',
                                     'white-collar','unemployed',
                                     'unemployed'])
```

```
In [11]: #Check marital status
df['marital'].unique()
```

```
Out[11]: array(['married', 'single', 'divorced', 'unknown'], dtype=object)
```

```
In [12]: df['marital']=df['marital'].replace(['divorced','unknown'],['single','marital_unknown'])
df['marital'].unique()
```

```
Out[12]: array(['married', 'single', 'marital_unknown'], dtype=object)
```

```
In [13]: #Check unique entries in education column
df['education'].unique()
```

```
Out[13]: array(['basic.4y', 'unknown', 'university.degree', 'high.school',
               'basic.9y', 'professional.course', 'basic.6y', 'illiterate'],
              dtype=object)
```

```
In [14]: #Replace
#basic.4y. basic.6y,basic.9y with basic
#unknown with edu_unknown
df['education']=df['education'].replace(['basic.4y','basic.6y','basic.9y'],
                                       ['basic','basic','basic'])
df['education']=df['education'].replace(['unknown'],
                                       ['edu_unknown'])
```

```
In [15]: #Check unique entries in education column
df['education'].unique()
```

```
Out[15]: array(['basic', 'edu_unknown', 'university.degree', 'high.school',
               'professional.course', 'illiterate'], dtype=object)
```

```
In [16]: df['default'].unique()
df['default']=df['default'].replace(['unknown','yes','no'],
                                   ['default_unknown','default_yes','default_no'])
```

```
In [17]: df['default'].unique()
```

```
Out[17]: array(['default_unknown', 'default_no', 'default_yes'], dtype=object)
```

```
In [18]: df['housing'].unique()
```

```
Out[18]: array(['yes', 'no', 'unknown'], dtype=object)
```

```
In [19]: df['housing']=df['housing'].replace(['unknown','yes','no'],
                                           ['housing_unknown','housing_yes','housing_no'])
df['housing'].unique()
```

```
Out[19]: array(['housing_yes', 'housing_no', 'housing_unknown'], dtype=object)
```

```
In [20]: df['loan'].unique()
df['loan']=df['loan'].replace(['unknown','yes','no'],
                             ['loan_unknown','loan_yes','loan_no'])
df['loan'].unique()
```

```
Out[20]: array(['loan_no', 'loan_yes', 'loan_unknown'], dtype=object)
```

```
In [21]: df['contact'].unique()
```

```
Out[21]: array(['cellular', 'telephone'], dtype=object)
```


```
In [22]: df['poutcome'].unique()
```

```
Out[22]: array(['nonexistent', 'success', 'failure'], dtype=object)
```


```
In [23]: #Check missing values
df.isnull().sum()
```

```
Out[23]: age                0
job                0
marital           0
education         0
default           0
housing           0
loan              0
contact           0
month             0
day_of_week       0
duration          0
campaign          0
pdays            0
previous          0
poutcome          0
emp_var_rate      0
cons_price_idx    0
cons_conf_idx     0
euribor3m         0
nr_employed       0
Invested          0
dtype:int64
```


```
In [24]: #Check outliers
plt.boxplot(df['age']) #has outlier
plt.show()
```



```
In [25]: #Check outliers
plt.boxplot(df['duration']) #has outlier
plt.show()
```



```
In [26]: #Check outliers
plt.boxplot(df['pdays']) #has outlier
plt.show()
```



```
In [30]: #remove outliers
#user defined function for outlier treatment
def rm_out(d,c):
    #find q1 and q3
    q1=d[c].quantile(0.25)
    q3=d[c].quantile(0.75)

    #find iqr
    iqr=q3-q1


    #upper bound and lower bound
    ub=q3+1.5*iqr
    lb=q1-1.5*iqr

    output_data=d.loc[(d[c]>lb) & (d[c]<ub)]

    return output_data
```


```
In [35]: #Outlier treatment: age column
df=rm_out(df,'age')
plt.boxplot(df['age'])
```

```
Out[35]: {'whiskers': [<matplotlib.lines.Line2D at 0x22dbe480610>,
<matplotlib.lines.Line2D at 0x22dbe4808e0>],
'caps': [<matplotlib.lines.Line2D at 0x22dbe480bb0>,
<matplotlib.lines.Line2D at 0x22dbe480e80>],
'boxes': [<matplotlib.lines.Line2D at 0x22dbe480340>],
'medians': [<matplotlib.lines.Line2D at 0x22dbe48d190>],
'fliers': [<matplotlib.lines.Line2D at 0x22dbe48d460>],
'means': []}
```



```
In [45]: #Outlier treatment: duration column
df=rm_out(df,'duration')
plt.boxplot(df['duration'])
```

```
Out[45]: {'whiskers': [<matplotlib.lines.Line2D at 0x22dbe65ad00>,
<matplotlib.lines.Line2D at 0x22dbe65afd0>],
'caps': [<matplotlib.lines.Line2D at 0x22dbe6692e0>,
<matplotlib.lines.Line2D at 0x22dbe6695b0>],
'boxes': [<matplotlib.lines.Line2D at 0x22dbe65aa30>],
'medians': [<matplotlib.lines.Line2D at 0x22dbe669880>],
'fliers': [<matplotlib.lines.Line2D at 0x22dbe669b50>],
'means': []}
```



```
In [48]: #Replace pdays 999 with NaN (missing values)
df['pdays']=df['pdays'].replace(999,np.nan)
```

```
In [49]: #Replace pdays NaN (missing values) with median
df['pdays']=df['pdays'].fillna(df['pdays'].median())
```

```
In [50]: df.columns
```

```
Out[50]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
'previous', 'poutcome', 'emp_var_rate', 'cons_price_idx',
'cons_conf_idx', 'euribor3m', 'nr_employed', 'Invested'],
dtype='object')
```

```
In [51]: # split data into x and y
x = df.loc[:, ['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
'previous', 'poutcome', 'emp_var_rate', 'cons_price_idx',
'cons_conf_idx', 'euribor3m', 'nr_employed']]

y = df['Invested']
```

```
In [52]: # create dummy variables for categorical variables

# subset all categorical variables
df_categorical = x.select_dtypes(include=['object'])
df_categorical.head()
```

Out[52]:

	job	marital	education	default	housing	loan	contact	month	day_of_week	poutcome
0	blue-collar	married	basic	default_unknown	housing_yes	loan_no	cellular	aug	thu	nonexistent
1	blue-collar	married	edu_unknown	default_no	housing_no	loan_no	cellular	nov	fri	nonexistent
2	white-collar	single	university.degree	default_no	housing_yes	loan_no	cellular	jun	thu	success
3	white-collar	married	high.school	default_no	housing_no	loan_no	cellular	apr	fri	nonexistent
4	retired	married	basic	default_no	housing_yes	loan_no	cellular	aug	fri	success

In [53]:

```
# convert into dummies
df_dummies = pd.get_dummies(df_categorical)
df_dummies.head()
```

Out[53]:

	job_blue-collar	job_retired	job_unemployed	job_white-collar	marital_marital_unknown	marital_married	marital_single	education_basic	education_ed
0	1	0	0	0		0	1	0	1
1	1	0	0	0		0	1	0	0
2	0	0	0	1		0	0	1	0
3	0	0	0	1		0	1	0	0
4	0	1	0	0		0	1	0	1

5 rows × 42 columns

In [54]:

```
# drop categorical from x
x = x.drop(list(df_categorical.columns), axis=1)
```

In [55]:

```
# concat dummy variables with x
x = pd.concat([x, df_dummies], axis=1)
x
```

Out[55]:

	age	duration	campaign	pdays	previous	emp_var_rate	cons_price_idx	cons_conf_idx	euribor3m	nr_employed	...	month_oct	mo
0	44	210	1	6.0	0	1.4	93.444	-36.1	4.963	5228.1	...	0	
1	53	138	1	6.0	0	-0.1	93.200	-42.0	4.021	5195.8	...	0	
2	28	339	3	6.0	2	-1.7	94.055	-39.8	0.729	4991.6	...	0	
3	39	185	2	6.0	0	-1.8	93.075	-47.1	1.405	5099.1	...	0	
4	55	137	1	3.0	1	-2.9	92.201	-31.4	0.869	5076.2	...	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
41183	59	222	1	6.0	0	1.4	94.465	-41.8	4.866	5228.1	...	0	
41184	31	196	2	6.0	0	1.1	93.994	-36.4	4.860	5191.0	...	0	
41185	42	62	3	6.0	0	1.1	93.994	-36.4	4.857	5191.0	...	0	
41186	48	200	2	6.0	0	-3.4	92.431	-26.9	0.742	5017.5	...	1	
41187	25	112	4	6.0	0	1.1	93.994	-36.4	4.859	5191.0	...	0	

35688 rows × 52 columns

In [56]:

```
# split data into train and test
from sklearn.model_selection import train_test_split

xtrain,xtest, ytrain, ytest = train_test_split(x, y,train_size=0.7,test_size = 0.3,
                                              random_state=0)
```

In [57]:

```
#----- Model Development -----
```

## ----- Logistic Regression -----

In [58]:

```
#Import Logistic Regression Library
from sklearn.linear_model import LogisticRegression
```

In [59]:

```
#create model object
logreg=LogisticRegression()
```

In [60]:

```
#Fit the model using training data
train_model_fit=logreg.fit(xtrain,ytrain)
```

In [61]:

```
#check the accuracy of training model
logreg.score(xtrain,ytrain)
```

Out[61]: 0.9465994155558224

```
In [62]: #predict y using test data
y_pred=logreg.predict(xtest)
```

```
In [63]: #check the accuracy of test model
logreg.score(xtest,ytest)
```

Out[63]: 0.9518072289156626

```
In [64]: from sklearn.metrics import accuracy_score, f1_score,recall_score,precision_score, confusion_matrix
acc = accuracy_score(ytest, y_pred)
prec = precision_score(ytest, y_pred)
rec = recall_score(ytest, y_pred)
f1 = f1_score(ytest,y_pred)
```

```
print('Accuracy:', acc,'\nPrecision:', prec,'\nRecall:',rec, '\nF1 Score:',f1)
```

```
results = pd.DataFrame(['Logistic Regression', acc,prec,rec,f1]),
                      columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results
```

Accuracy: 0.9518072289156626  
Precision: 0.6782178217821783  
Recall: 0.41515151515151516  
F1 Score: 0.5150375939849625

```
Out[64]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Logistic Regression	0.951807	0.678218	0.415152	0.515038

```
In [65]: prec = precision_score(ytest, y_pred)
prec
```

Out[65]: 0.6782178217821783

```
In [66]: #import confusion matrix library
from sklearn.metrics import confusion_matrix
```

```
In [67]: #create confusion matrix
confusion_matrix(ytest,y_pred)
```

```
Out[67]: array([[9917, 130],
 [ 386, 274]], dtype=int64)
```

```
In [68]: #Precision: Out of all predicted positives, how many are really positive
#Precision=TP/(TP+FP)
274/(274+130)
```

Out[68]: 0.6782178217821783

```
In [69]: #Accuracy =(TP+TN)/(TP+TN+FN+FP)
(9917+274)/(9917+274+386+130)
```

Out[69]: 0.9518072289156626

```
In [70]: #TPR or sensitivity or recall TPR=TP/TP+FN
274/(386+274) #sensitivity
```

Out[70]: 0.41515151515151516

```
In [71]: #TNR or Specificity or selectivity TNR =TN/TN+FP
9917/(9917+ 130)
```

Out[71]: 0.9870608141733851

## -----KNN -----

```
In [72]: #import knn library from sklearn
from sklearn.neighbors import KNeighborsClassifier
```

```
In [73]: #Create model object
knn=KNeighborsClassifier(n_neighbors=5, metric='euclidean')
```

```
In [74]: #fit the training model
train_model_fit=knn.fit(xtrain,ytrain)
```

```
In [75]: #test the accuracy of training model
knn.score(xtrain,ytrain)
```

Out[75]: 0.9559265041431488

```
In [76]: #test the model using xtest
y_pred=knn.predict(xtest)
```

```
In [77]: #Check accuracy of test model
knn.score(xtest,ytest)
```

Out[77]: 0.9462034183244606

```
In [78]: #import confusion matrix from sklearn
from sklearn.metrics import confusion_matrix
```

```
In [79]: confusion_matrix(ytest,y_pred)
```

Out[79]: array([[9853, 194],  
[ 382, 278]], dtype=int64)

```
In [80]: #TPR or Sensitivity or recall
278/(382+ 278)
```

Out[80]: 0.4212121212121212

```
In [81]: #TNR or Specificity or Selectivity
9852/(9852+ 195)
```

Out[81]: 0.9805912212600776

```
In [82]: #Calculate Recall, Precision, Sensitivity
knn_acc=(9854+279)/(9854+279+193+381)
knn_rec=279/(279+381)
knn_prec=(279/(279+193))
knn_spec=9854/(9854+193)
knn_f1=2*(knn_prec*knn_rec)/(knn_prec + knn_rec)
```

```
In [83]: results = pd.DataFrame([[ 'Log Reg', acc,prec,rec,f1],
                                [ 'KNN', knn_acc,knn_prec,knn_rec,knn_f1]],
                                columns=[ 'Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results
```

Out[83]:

	Model	Accuracy	Precision	Recall	F1 Score
0	Log Reg	0.951807	0.678218	0.415152	0.515038
1	KNN	0.946390	0.591102	0.422727	0.492933

In [ ]:

```
In [84]: # predict probabilities
pred_prob1 = logreg.predict_proba(xtest)
pred_prob2 = knn.predict_proba(xtest)
```

```
In [85]: from sklearn.metrics import roc_curve

# roc curve for models
fpr1, tpr1, thresh1 = roc_curve(ytest, pred_prob1[:,1], pos_label=1)
fpr2, tpr2, thresh2 = roc_curve(ytest, pred_prob2[:,1], pos_label=1)

# roc curve for tpr = fpr
random_probs = [0 for i in range(len(ytest))]
p_fpr, p_tpr, _ = roc_curve(ytest, random_probs, pos_label=1)
```

```
In [86]: from sklearn.metrics import roc_auc_score

# auc scores
auc_score1 = roc_auc_score(ytest, pred_prob1[:,1])
auc_score2 = roc_auc_score(ytest, pred_prob2[:,1])

print(auc_score1, auc_score2)

0.937673992839714 0.8623289478843376
```


```
In [87]: # matplotlib
import matplotlib.pyplot as plt
plt.style.use('seaborn')

# plot roc curves
plt.plot(fpr1, tpr1, linestyle='--',color='orange', label='Logistic Regression')
plt.plot(fpr2, tpr2, linestyle='--',color='green', label='KNN')
plt.plot(p_fpr, p_tpr, linestyle='--', color='blue')
# title
plt.title('ROC curve')
# x label
plt.xlabel('False Positive Rate')
```



```
# y label
plt.ylabel('True Positive rate')

plt.legend(loc='best')
plt.savefig('ROC')
plt.show()
```



## ----- Decision Tree -----

```
In [105... #import Decision Tree classifier library from sklearn
from sklearn.tree import DecisionTreeClassifier
```

```
In [106... #Create model object
dtree=DecisionTreeClassifier(max_depth=5)
```

```
In [107... #fit the training model
model=dtree.fit(xtrain,ytrain)
```

```
In [108... #test the accuracy of training model
dtree.score(xtrain,ytrain)
```

```
Out[108]: 0.9525639486009367
```

```
In [109... #test the model using xtest
y_pred=dtree.predict(xtest)
```

```
In [110... #Check accuracy of test model
dtree.score(xtest,ytest)
```

```
Out[110]: 0.9537685626225834
```

```
In [111... #import confusion matrix from sklearn
from sklearn.metrics import confusion_matrix
```

```
In [112... confusion_matrix(ytest,y_pred)
```

```
Out[112]: array([[9813, 234],
               [ 261, 399]], dtype=int64)
```

```
In [113... #TPR or Sensitivity or recall
399/(261+399)
```

```
Out[113]: 0.6045454545454545
```

```
In [114... #TNR Specificity or selectivity
9813/(9813+234)
```

```
Out[114]: 0.9767094655120931
```

```
In [115... from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score, confusion_matrix
dt_acc = accuracy_score(ytest, y_pred)
dt_prec = precision_score(ytest, y_pred)
dt_rec = recall_score(ytest, y_pred)
dt_f1 = f1_score(ytest, y_pred)

results = pd.DataFrame(['Log Reg', acc, prec, rec, f1],
                        ['KNN', knn_acc, knn_prec, knn_rec, knn_f1],
                        ['Naive Bayes', nb_acc, nb_prec, nb_rec, nb_f1],
                        ['D-Tree', dt_acc, dt_prec, dt_rec, dt_f1])
```

```
columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results
```

```
Out[115]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Log Reg	0.951807	0.678218	0.415152	0.515038
1	KNN	0.946390	0.591102	0.422727	0.492933
2	Naive Bayes	0.881760	0.312268	0.763636	0.443272
3	D-Tree	0.953769	0.630332	0.604545	0.617169

## ----- Random FOrest -----

```
In [ ]:
```

```
In [116.. #import Random Forest classifier library from sklearn
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import StratifiedKFold
from sklearn.feature_selection import RFECV
```

```
In [137.. #Create Model Object for Random Forest
rfc = RandomForestClassifier(random_state=101)

#Recursive Feature Elimination with Cross-Validation
rfecv = RFECV(estimator=rfc, step=1, cv=StratifiedKFold(7),
               scoring='accuracy')
rfecv
```

```
Out[137]: RFECV(cv=StratifiedKFold(n_splits=7, random_state=None, shuffle=False),
               estimator=RandomForestClassifier(random_state=101), scoring='accuracy')
```

```
In [138.. #fit the training model
model=rfc.fit(xtrain,ytrain)
```

```
In [139.. #test the accuracy of training model
rfc.score(xtrain,ytrain)
```

```
Out[139]: 1.0
```

```
In [140.. #test the model using xtest
y_pred=rfc.predict(xtest)
```

```
In [141.. #Check accuracy ot test model
rfc.score(xtest,ytest)
```

```
Out[141]: 0.9545157373680769
```

```
In [142.. #import confusion matrix from sklearn
from sklearn.metrics import confusion_matrix
```

```
In [143.. confusion_matrix(ytest,y_pred)
```

```
Out[143]: array([[9887, 160],
                [ 327, 333]], dtype=int64)
```

```
In [144.. #TPR or Sensitivity or recall
334/(334+326)
```

```
Out[144]: 0.5060606060606061
```

```
In [145.. #TNR Spcificity or selectivity
9885/(9813+162)
```

```
Out[145]: 0.9909774436090225
```

```
In [146.. from sklearn.metrics import accuracy_score, f1_score,recall_score,precision_score, confusion_matrix
rf_acc = accuracy_score(ytest, y_pred)
rf_prec = precision_score(ytest, y_pred)
rf_rec = recall_score(ytest, y_pred)
rf_f1 = f1_score(ytest,y_pred)

results = pd.DataFrame(['Log Reg', acc,prec,rec,f1],
                        ['KNN', knn_acc,knn_prec,knn_rec,knn_f1],
                        ['Naive Bayes', nb_acc,nb_prec,nb_rec,nb_f1],
                        ['D-Tree', dt_acc,dt_prec,dt_rec,dt_f1],
                        ['Randm Forest', rf_acc,rf_prec,rf_rec,rf_f1]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
results
```

```
Out[146]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Log Reg	0.951807	0.678218	0.415152	0.515038
1	KNN	0.946390	0.591102	0.422727	0.492933
2	Naive Bayes	0.881760	0.312268	0.763636	0.443272
3	D-Tree	0.953769	0.630332	0.604545	0.617169
4	Randm Forest	0.954516	0.675456	0.504545	0.577624

## -----Support Vector Machine (SVM) -----

```
In [127... #Support Vector Machine (SVC: Support Vector Classifier)
from sklearn.svm import SVC
svm = SVC(kernel='linear')
```

```
In [128... #Fit the model
model=svm.fit(xtrain,ytrain)
```

```
In [129... #check accuracy of training model
svm.score(xtrain,ytrain)
```

```
Out[129]: 0.9360313838517273
```

```
In [130... #Predict the response from xtest
y_pred=svm.predict(xtest)
```

```
In [131... #Check accuracy of test model
svm.score(xtest,ytest)
```

```
Out[131]: 0.9343420192397497
```

```
In [132... #import confusion matrix from sklearn
from sklearn.metrics import confusion_matrix
```

```
In [133... confusion_matrix(ytest,y_pred)
```

```
Out[133]: array([[9698, 349],
 [ 354, 306]], dtype=int64)
```

```
In [134... #Sensitivity
306/(354+306)
```

```
Out[134]: 0.4636363636363636
```

```
In [135... #Specificity
9718/(9718+329)
```

```
Out[135]: 0.9672539066387976
```

```
In [136... from sklearn.metrics import accuracy_score, f1_score,recall_score,precision_score, confusion_matrix
svm_acc = accuracy_score(ytest, y_pred)
svm_prec = precision_score(ytest, y_pred)
svm_rec = recall_score(ytest, y_pred)
svm_f1 = f1_score(ytest,y_pred)

results = pd.DataFrame([['Log Reg', acc,prec,rec,f1],
                        ['KNN', knn_acc,knn_prec,knn_rec,knn_f1],
                        ['Naive Bayes', nb_acc,nb_prec,nb_rec,nb_f1],
                        ['D-Tree', dt_acc,dt_prec,dt_rec,dt_f1],
                        ['Randm Forest', rf_acc,rf_prec,rf_rec,rf_f1],
                        ['SVM', svm_acc,svm_prec,svm_rec,svm_f1]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall','F1 Score'])
results
```

```
Out[136]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Log Reg	0.951807	0.678218	0.415152	0.515038
1	KNN	0.946390	0.591102	0.422727	0.492933
2	Naive Bayes	0.881760	0.312268	0.763636	0.443272
3	D-Tree	0.953769	0.630332	0.604545	0.617169
4	Randm Forest	0.954516	0.675456	0.504545	0.577624
5	SVM	0.934342	0.467176	0.463636	0.465399

## ----- Naive Bayes Classifier -----

```
In [99]: #Import Gaussian Naive Bayes model
from sklearn.naive_bayes import GaussianNB

#Create a Gaussian Classifier
gnb = GaussianNB()

# Train the model using the training sets
gnb.fit(xtrain,ytrain)

#Predict Output
y_pred = gnb.predict(xtest)
```

```
In [100]: #Check accuracy of test model

#use model.score(xtest,ytest)
gnb.score(xtest,ytest)
#or
# Model Accuracy, how often is the classifier correct?
from sklearn import metrics
print("Accuracy:",metrics.accuracy_score(ytest, y_pred))
```

Accuracy: 0.8817595965256374

```
In [101]: #TPR
from sklearn.metrics import confusion_matrix

confusion_matrix(ytest,y_pred)
#TPR=0.85
```

```
Out[101]: array([[ 8937, 1110],
               [ 156,  504]], dtype=int64)
```

```
In [102]: #sensitivity
504/(156+504)
```

```
Out[102]: 0.7636363636363637
```

```
In [103]: #specificity
8937/(8937+1110)
```

```
Out[103]: 0.889519259480442
```

```
In [104]: from sklearn.metrics import accuracy_score, f1_score, recall_score, precision_score, confusion_matrix
nb_acc = accuracy_score(ytest, y_pred)
nb_prec = precision_score(ytest, y_pred)
nb_rec = recall_score(ytest, y_pred)
nb_f1 = f1_score(ytest, y_pred)

#print('Accuracy:', nb_acc, '\nPrecision:', nb_prec, '\nRecall:', nb_rec, '\nF1 Score:', nb_f1)

results = pd.DataFrame([['Log Reg', acc, prec, rec, f1],
                        ['KNN', knn_acc, knn_prec, knn_rec, knn_f1],
                        ['Naive Bayes', nb_acc, nb_prec, nb_rec, nb_f1]],
                        columns=['Model', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])
results
```

```
Out[104]:
```

	Model	Accuracy	Precision	Recall	F1 Score
0	Log Reg	0.951807	0.678218	0.415152	0.515038
1	KNN	0.946390	0.591102	0.422727	0.492933
2	Naive Bayes	0.881760	0.312268	0.763636	0.443272

```
In [ ]:
```

```
In [ ]:
```